

Assignment 4

Kate Lamoreaux

2023-10-25

Link to github repository: https://github.com/katelmrx/SURV727_Assignment4

In this notebook we will use Google BigQuery, “Google’s fully managed, petabyte scale, low cost analytics data warehouse”. Some instruction on how to connect to Google BigQuery can be found here: <https://db.rstudio.com/databases/big-query/>.

You will need to set up a Google account with a project to be able to use this service. We will be using a public dataset that comes with 1 TB/mo of free processing on Google BigQuery. As long as you do not repeat the work in this notebook constantly, you should be fine with just the free tier.

Go to <https://console.cloud.google.com> and make sure you are logged in a non-university Google account. **This may not work on a university G Suite account because of restrictions on those accounts.** Create a new project by navigating to the dropdown menu at the top (it might say “Select a project”) and selecting “New Project” in the window that pops up. Name it something useful.

```
project <- "surv727-fcdd"
```

After you have initialized a project, paste your project ID into the following chunk.

```
con <- dbConnect(  
  bigrquery::bigquery(),  
  project = "bigquery-public-data",  
  dataset = "chicago_crime",  
  billing = project  
)  
con
```

We will connect to a public database, the Chicago crime database, which has data on crime in Chicago.

```
## <BigQueryConnection>  
##   Dataset: bigquery-public-data.chicago_crime  
##   Billing: surv727-fcdd
```

We can look at the available tables in this database using `dbListTables`.

Note: When you run this code, you will be sent to a browser and have to give Google permissions to Tidyverse API Packages. **Make sure you select all to give access or else your code will not run.**

```
dbListTables(con)
```

```
## ! Using an auto-discovered, cached token.  
##   To suppress this message, modify your code or options to clearly consent to  
##   the use of a cached token.
```

```
## See gargle's "Non-interactive auth" vignette for more details:
## <https://gargle.r-lib.org/articles/non-interactive-auth.html>
## i The bigrquery package is using a cached token for 'katelmr@gmail.com'.
## [1] "crime"
```

Information on the 'crime' table can be found here:

<https://cloud.google.com/bigquery/public-data/chicago-crime-data>

```
SELECT COUNT(*) as COUNT
FROM crime
WHERE year = 2016
LIMIT 10;
```

Write a first query that counts the number of rows of the 'crime' table in the year 2016. Use code chunks with {sql connection = con} in order to write SQL code within the document.

Table 1: 1 records

COUNT
269854

According to our first query, there are 269,854 rows in the crime table for the year 2016.

```
SELECT primary_type, COUNT(*) as COUNT
FROM crime
WHERE year= 2016 and arrest = TRUE
GROUP BY primary_type
ORDER BY COUNT(*) DESC
LIMIT 10;
```

Next, count the number of arrests grouped by primary_type in 2016. Note that is a somewhat similar task as above, with some adjustments on which rows should be considered. Sort the results, i.e. list the number of arrests in a descending order.

Table 2: Displaying records 1 - 10

primary_type	COUNT
NARCOTICS	13327
BATTERY	10332
THEFT	6522
CRIMINAL TRESPASS	3724
ASSAULT	3492
OTHER OFFENSE	3415
WEAPONS VIOLATION	2511
CRIMINAL DAMAGE	1669
PUBLIC PEACE VIOLATION	1116
MOTOR VEHICLE THEFT	1097

Arrests for narcotics, battery, and theft were the most common in 2016.

We can also use the `date` for grouping. Count the number of arrests grouped by hour of the day in 2016. You can extract the latter information from `date` via `EXTRACT(HOUR FROM date)`. Which time of the day is associated with the most arrests?

```
SELECT EXTRACT(HOUR FROM date) AS hour_of_day, COUNT(*) as COUNT
FROM crime
WHERE year= 2016 and arrest = TRUE
GROUP BY hour_of_day
ORDER BY COUNT(*) DESC
LIMIT 12;
```

Table 3: Displaying records 1 - 10

hour_of_day	COUNT
10	5306
11	5200
12	4944
7	4900
8	4735
9	4675
1	4288
6	4261
2	4029
3	3750

Later hours of the data are associated with more arrests. Hour 10 is associated with the most arrests.

```
SELECT YEAR, COUNT(*) as COUNT
FROM crime
WHERE primary_type='HOMICIDE' and arrest = TRUE
GROUP BY year
ORDER BY COUNT(*) DESC
LIMIT 20;
```

Focus only on HOMICIDE and count the number of arrests for this incident type, grouped by year. List the results in descending order.

Table 4: Displaying records 1 - 10

YEAR	COUNT
2001	430
2002	424
2003	379
2020	341
2004	293
2008	286
2016	286
2005	281
2006	281
2022	280

Years 2001-2003 had the highest number of arrests for homicides.

```
SELECT YEAR, DISTRICT, COUNT(*) as COUNT
FROM crime
WHERE arrest = TRUE AND (YEAR = 2015 OR YEAR = 2016)
GROUP BY YEAR, DISTRICT
ORDER BY COUNT(*) DESC
LIMIT 20;
```

Find out which districts have the highest numbers of arrests in 2015 and 2016. That is, count the number of arrests in 2015 and 2016, grouped by year and district. List the results in descending order.

Table 5: Displaying records 1 - 10

YEAR	DISTRICT	COUNT
2015	11	8974
2016	11	6575
2015	7	5549
2015	15	4514
2015	6	4473
2015	25	4448
2015	4	4325
2015	8	4112
2016	7	3654
2015	10	3621

Districts 11 and 7, respectively, had the highest and second-highest numbers of arrests in both 2015 and 2016.

```
dbquery1 <- "SELECT primary_type, COUNT(*) as COUNT
FROM crime
WHERE (arrest = TRUE AND year = 2016) AND district = 11
GROUP BY primary_type
ORDER BY COUNT(*) DESC
LIMIT 10"
```

Let's switch to writing queries from within R via the DBI package. Create a query object that counts the number of arrests grouped by `primary_type` of district 11 in year 2016. The results should be displayed in descending order. The above code assigns `dbquery1` as a query object that counts the number of arrests in descending order grouped by `primary_type` within District 11 in the year 2016.

```
dbGetQuery(con, dbquery1)
```

Execute the query.

```
## # A tibble: 10 x 2
##   primary_type      COUNT
##   <chr>          <int>
## 1 NARCOTICS      3634
## 2 BATTERY        635
```

```
## 3 PROSTITUTION 511
## 4 WEAPONS VIOLATION 303
## 5 OTHER OFFENSE 255
## 6 ASSAULT 206
## 7 CRIMINAL TRESPASS 205
## 8 PUBLIC PEACE VIOLATION 135
## 9 INTERFERENCE WITH PUBLIC OFFICER 119
## 10 CRIMINAL DAMAGE 106
```

The above code uses the function `dbGetQuery()` to execute query `dbquery1`. 2016 arrests in District 11 were mostly for narcotics, battery, and prostitution.

```
crimetibble <- tbl(con, 'crime')
```

Try to write the very same query, now using the `dbplyr` package. For this, you need to first map the crime table to a tibble object in R.

```
## Warning: <BigQueryConnection> uses an old dbplyr interface
## i Please install a newer version of the package or contact the maintainer
## This warning is displayed once every 8 hours.
```

The above code maps the crime table to a tibble object in R, titled `crimetibble`.

```
crimetibble %>%
  filter(arrest == TRUE, year == 2016, district == 11) %>%
  group_by(primary_type) %>%
  summarize(count = n()) %>%
  arrange(desc(count))
```

Again, count the number of arrests grouped by `primary_type` of district 11 in year 2016, now using `dplyr` syntax.

```
## # Source:      SQL [?? x 2]
## # Database:    BigQueryConnection
## # Ordered by: desc(count)
##   primary_type      count
##   <chr>             <int>
## 1 NARCOTICS         3634
## 2 BATTERY           635
## 3 PROSTITUTION      511
## 4 WEAPONS VIOLATION 303
## 5 OTHER OFFENSE     255
## 6 ASSAULT           206
## 7 CRIMINAL TRESPASS 205
## 8 PUBLIC PEACE VIOLATION 135
## 9 INTERFERENCE WITH PUBLIC OFFICER 119
## 10 CRIMINAL DAMAGE  106
## # i more rows
```

The above code counts the number of arrests in descending order grouped by `primary_type` within District 11 in the year 2016, using `dplyr` syntax. Again, we see the same numbers as we did in DBI query `dbquery1`.

```
crimetibble %>%
  filter(arrest == TRUE, district == 11) %>%
```

```
group_by(year, primary_type) %>%
summarize(count = n()) %>%
arrange(desc(year))
```

Count the number of arrests grouped by `primary_type` and `year`, still only for district 11. Arrange the result by `year`.

```
## `summarise()` has grouped output by "year". You can override using the
## `.groups` argument.
```

```
## # Source:      SQL [?? x 3]
## # Database:    BigQueryConnection
## # Groups:      year
## # Ordered by:  desc(year)
##   year primary_type      count
##   <int> <chr>             <int>
## 1  2023 WEAPONS VIOLATION    437
## 2  2023 PROSTITUTION        143
## 3  2023 CRIMINAL TRESPASS    51
## 4  2023 HOMICIDE            17
## 5  2023 OFFENSE INVOLVING CHILDREN 10
## 6  2023 SEX OFFENSE          5
## 7  2023 ASSAULT             114
## 8  2023 ROBBERY              28
## 9  2023 OTHER OFFENSE       189
## 10 2023 LIQUOR LAW VIOLATION    5
## # i more rows
```

This code counts the number of arrests grouped by `primary_type` and `year`, still only for district 11. With the results arranged by `year`, however, the count of arrests is no longer in descending order.

```
arrestsbytypeandyear_chidistrict11 <- crimetibble %>%
  filter(arrest == TRUE, district == 11) %>%
  group_by(year, primary_type) %>%
  summarize(count = n()) %>%
  arrange(desc(year))
```

Assign the results of the query above to a local R object. The above code assigns the results of our prior SQL query using `dplyr` syntax to a local R object titled, `arrestsbytypeandyear_chidistrict11`.

```
head(arrestsbytypeandyear_chidistrict11, 10)
```

Confirm that you pulled the data to the local environment by displaying the first ten rows of the saved data set.

```
## `summarise()` has grouped output by "year". You can override using the
## `.groups` argument.

## # Source:      SQL [10 x 3]
## # Database:    BigQueryConnection
## # Groups:      year
## # Ordered by:  desc(year)
##   year primary_type      count
##   <int> <chr>             <int>
```

```
## 1 2023 INTERFERENCE WITH PUBLIC OFFICER      43
## 2 2023 OFFENSE INVOLVING CHILDREN             10
## 3 2023 PROSTITUTION                          143
## 4 2023 THEFT                                  27
## 5 2023 CRIMINAL DAMAGE                       56
## 6 2023 CONCEALED CARRY LICENSE VIOLATION      5
## 7 2023 NARCOTICS                             1415
## 8 2023 WEAPONS VIOLATION                     437
## 9 2023 STALKING                              3
## 10 2023 PUBLIC PEACE VIOLATION                28
```

The above code confirms I pulled the data into my local environment by printing the first ten rows of my new R object, `arrestsbytypeandyear_chidistrict11`.

```
#code to close database connection
dbDisconnect(con)
```

Close the connection.